

```
#import necessary libraries
import pandas as pd
import numpy as np
import pickle
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import sklearn

from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.model_selection import RandomizedSearchCV
import imblearn
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, f1_score

#import dataset
data = pd.read_csv(r"Churn_Modelling.csv")
data
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
1	2	Hargrave	619	France	Female	42	2
2	3	Onio	608	Spain	Female	41	1
3	4	Boni	502	France	Female	42	8
4	5	Mitchell	699	France	Female	39	1
...
9995	9996	Obijaku	850	Spain	Female	43	2
9996	9997	Johnstone	771	France	Male	39	5
9997	9998	Liu	516	France	Male	35	10
9998	9999	Sabbatini	709	France	Female	36	7
9999	10000	Walker	772	Germany	Male	42	3
10000	15628319		792	France	Female	28	4

10000 rows × 14 columns

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   RowNumber        10000 non-null   int64  
 1   CustomerId       10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   int64  
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   int64  
 7   Tenure            10000 non-null   int64  
 8   Balance           10000 non-null   float64 
 9   NumOfProducts     10000 non-null   int64  
 10  HasCrCard        10000 non-null   int64  
 11  IsActiveMember    10000 non-null   int64  
 12  EstimatedSalary   10000 non-null   float64 
 13 Exited             10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
#checking for null values
data.CustomerId=pd.to_numeric(data.CustomerId, errors='coerce')
data.isnull().any()
```

```
RowNumber      False
CustomerId     False
Surname       False
CreditScore    False
Geography     False
Gender        False
Age           False
Tenure         False
Balance        False
NumOfProducts  False
HasCrCard     False
IsActiveMember False
EstimatedSalary False
Exited        False
dtype: bool
```

```
data["CustomerId"].fillna(data["CustomerId"].median(), inplace=True)
data.isnull().sum()
```

```
RowNumber      0
CustomerId     0
Surname       0
CreditScore    0
Geography     0
Gender        0
Age           0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard     0
```

Saving...

```
Exited        0
dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data["RowNumber"] = le.fit_transform(data["RowNumber"])
data["CustomerId"] = le.fit_transform(data["CustomerId"])
data["Surname"] = le.fit_transform(data["Surname"])
data["CreditScore"] = le.fit_transform(data["CreditScore"])
data["Geography"] = le.fit_transform(data["Geography"])
data["Gender"] = le.fit_transform(data["Gender"])
data["Age"] = le.fit_transform(data["Age"])
data["Tenure"] = le.fit_transform(data["Tenure"])
data["Balance"] = le.fit_transform(data["Balance"])
data["NumOfProducts"] = le.fit_transform(data["NumOfProducts"])
data["HasCrCard"] = le.fit_transform(data["HasCrCard"])
data["IsActiveMember"] = le.fit_transform(data["IsActiveMember"])
data["EstimatedSalary"] = le.fit_transform(data["EstimatedSalary"])
data["Exited"] = le.fit_transform(data["Exited"])
```

```
data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	0	2736	1115	228	0	0	24	2						
1	1	3258	1177	217	2	0	23	1						
2	2	2104	2040	111	0	0	24	8						
3	3	5435	289	308	0	0	21	1						
4	4	6899	1822	459	2	0	25	2						

```
x=data.iloc[:,0:14].values
y=data.iloc[:,13:14].values
```

```
x
```

```
array([[ 0, 2736, 1115, ..., 1, 5068, 1],
       [ 1, 3258, 1177, ..., 1, 5639, 0],
       [ 2, 2104, 2040, ..., 0, 5707, 1],
       ...,
       [9997, 717, 1570, ..., 1, 2062, 1],
```

```
[9998, 4656, 2345, ...], 0, 4639, 1],
[9999, 2497, 2751, ...], 0, 1878, 0]])
```

y

```
array([[1],
       [0],
       [1],
       ...,
       [1],
       [1],
       [0]])
```

```
from sklearn.preprocessing import OneHotEncoder
one = OneHotEncoder()
a= one.fit_transform(x[:,0:1]).toarray()
b= one.fit_transform(x[:,1:2]).toarray()
c= one.fit_transform(x[:,2:3]).toarray()
d= one.fit_transform(x[:,3:4]).toarray()
e= one.fit_transform(x[:,4:5]).toarray()
f= one.fit_transform(x[:,5:6]).toarray()
g= one.fit_transform(x[:,6:7]).toarray()
h= one.fit_transform(x[:,7:8]).toarray()
i= one.fit_transform(x[:,8:9]).toarray()
j= one.fit_transform(x[:,9:10]).toarray()
x=np.delete(x, [0,1,2,3,4,5,6,7,8,9], axis=1)
x=np.concatenate((a,b,c,d,e,f,g,h,i,j,x), axis=1)
```

Saving... MOTE

```
smt = SMOTE()
x_resample,y_resample=smt.fit_resample (x,y)
```

x_resample

```
array([[1.00000000e+00, 0.00000000e+00, 0.00000000e+00, ...,
       1.00000000e+00, 5.06800000e+03, 1.00000000e+00],
       [0.00000000e+00, 1.00000000e+00, 0.00000000e+00, ...,
       1.00000000e+00, 5.63900000e+03, 0.00000000e+00],
       [0.00000000e+00, 0.00000000e+00, 1.00000000e+00, ...,
       0.00000000e+00, 5.70700000e+03, 1.00000000e+00],
       ...,
       [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ...,
       0.00000000e+00, 9.87599651e+03, 1.00000000e+00],
       [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ...,
       5.98737695e-01, 6.03421010e+03, 1.00000000e+00],
       [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ...,
       0.00000000e+00, 9.51229892e+03, 1.00000000e+00]])
```

y_resample

```
array([1, 0, 1, ..., 1, 1, 1])
```

data.describe()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gen
count	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000
mean	4999.50000	4999.50000	1507.774200	259.584600	0.746300	0.5451
std	2886.89568	2886.89568	846.204311	96.496107	0.827529	0.4971
min	0.00000	0.00000	0.00000	0.00000	0.00000	0.0001
25%	2499.75000	2499.75000	773.75000	193.00000	0.00000	0.0001
50%	4999.50000	4999.50000	1542.00000	261.00000	0.00000	1.0001
75%	7499.25000	7499.25000	2238.25000	327.00000	1.00000	1.0001
max	9999.00000	9999.00000	2931.00000	459.00000	2.00000	1.0001

```
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
```

```
sns.distplot(data["Balance"])
plt.subplot(1,2,2)
sns.distplot(data["Tenure"])
```

<ipython-input-21-40ad5a6b3819>:3: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

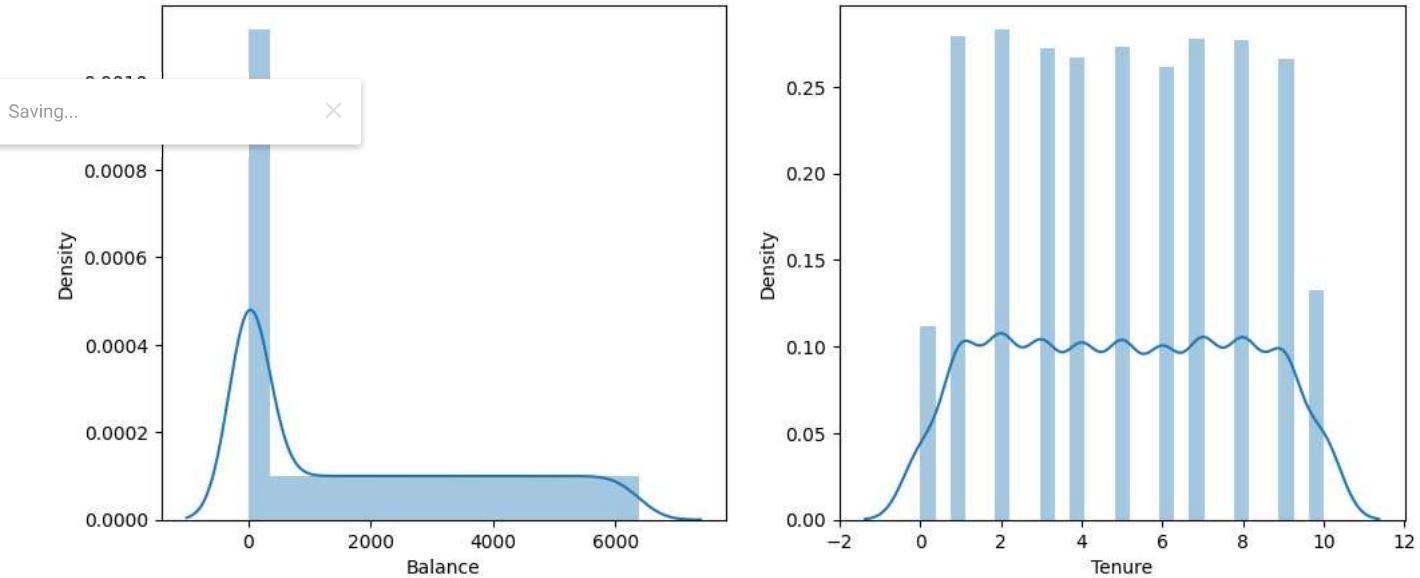
```
sns.distplot(data["Balance"])
<ipython-input-21-40ad5a6b3819>:5: UserWarning:  

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(data["Tenure"])
Axes: xlabel='Tenure', ylabel='Density'
```

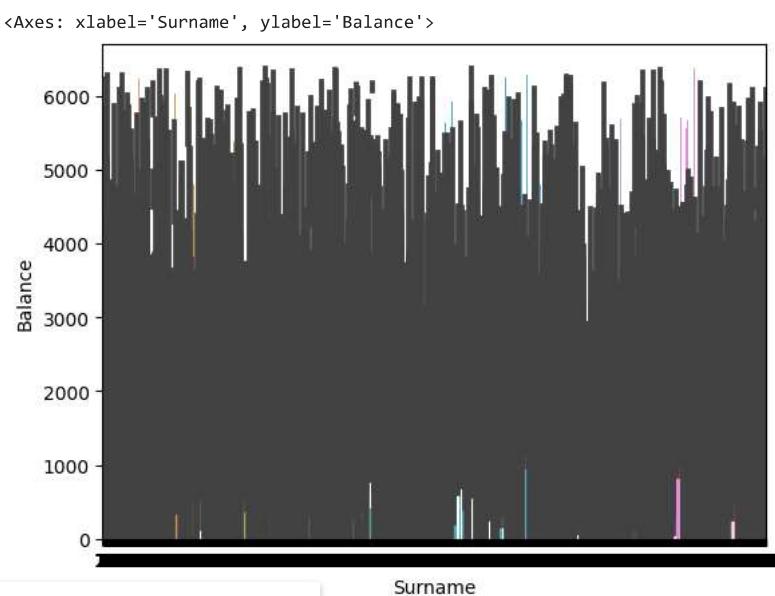


```
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
sns.countplot(data["Tenure"])
plt.subplot(1,2,1)
sns.countplot(data["Balance"])
```

```
<Axes: ylabel='count'>
```



```
sns.barplot(x="Surname", y="Balance", data=data)
```

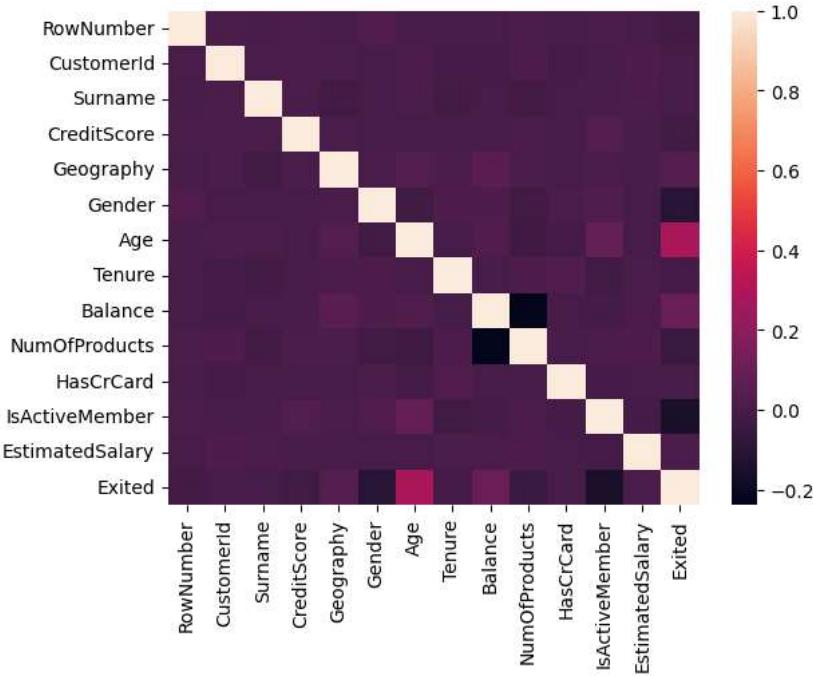


Saving...



```
sns.heatmap(data.corr(), annot=False)
```

```
<Axes: >
```



```
sns.pairplot(data=data, markers= ["^", "v"], palette= "inferno")
```

Saving...

4/12/23, 3:19 PM

Untitled1.ipynb - Colaboratory

```
/usr/local/lib/python3.9/dist-packages/seaborn/axisgrid.py:1609: UserWarning: Ignoring palette because no hue variable has been ass
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x_resample, y_resample, test_size = 0.2, random_state =0)
/usr/local/lib/python3.9/dist-packages/seaborn/axisgrid.py:1609: UserWarning: Ignoring palette because no hue variable has been ass
x_train.ndim

2
tunc(x=x, y=y, **kwargs)

x_test.ndim

2
tunc(x=x, y=y, **kwargs)

x_train=pd.DataFrame(x_train)
x_test=pd.DataFrame(x_test)
tunc(x=x, y=y, **kwargs)

x_train
```

	0	1	2	3	4	5	6	7	8	9	...	29858	29859	29860	29861	29862	29863	29864	29865	29866
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.000000	0.000000	0.000000	0.0	0.930159	0.069841	6636.768249
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.000000	0.000000	0.000000	0.0	0.000000	0.000000	7113.000000
Saving...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.043995	0.956005	0.000000	0.0	0.043995	1.000000	9852.296088
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.000000	0.000000	0.000000	0.0	1.000000	0.000000	5534.000000
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.000000	0.000000	0.000000	0.0	0.159173	0.000000	4274.522482
...	
12735	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.209732	0.790268	0.000000	0.0	0.209732	0.000000	6568.951340
12736	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.000000	1.000000	0.000000	0.0	1.000000	0.000000	9083.000000
12737	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.000000	1.000000	0.000000	0.0	1.000000	1.000000	7470.000000
12738	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.123083	0.000000	0.876917	0.0	0.123083	0.000000	391.107750
12739	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.000000	0.000000	0.000000	0.0	1.000000	0.000000	5952.000000

12740 rows × 29868 columns

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)
tunc(x=x, v=v, **kwargs)
x_train.shape

# importing and building the Decision tree model
def logreg(x_train,x_test,y_train,y_test):
    lr = LogisticRegression (random_state=0)
    lr.fit(x_train,y_train)
    y_lr_tr = lr.predict(x_train)
    print(accuracy_score(y_lr_tr,y_train))
    yPred_lr = lr.predict(x_test)
    print(accuracy_score (yPred_lr,y_test))
    print("****Logistic Regression****")
    print("Confusion_Matrix")
    print(confusion_matrix(y_test,yPred_lr))
    print("Classification Report")
    print(classification_report(y_test,yPred_lr))

# printing the train accuracy and test accuracy respectively
logreg(x_train,x_test,y_train,y_test)

# importing and building the Decision tree model
def decisionTree(x_train,x_test,y_train,y_test):
    dtc = DecisionTreeClassifier(criterion="entropy", random_state=0)
```

```

dtc.fit(x_train,y_train)
y_dt_tr = dtc.predict(x_train)
print(accuracy_score(y_dt_tr,y_train))
yPred_dt = dtc.predict(x_test)
print(accuracy_score(yPred_dt,y_test))
print("****Decision Tree****")
print("Confusion_Matrix")
print(confusion_matrix(y_test,yPred_dt))
print("Classification Report")
print(classification_report(y_test,yPred_dt))

#printing the train accuracy and test accuracy respectively
decisionTree(x_train,x_test,y_train,y_test)

#importing and building the random forest model
def RandomForest(x_tarin,x_test,y_train,y_test):
    rf = RandomForestClassifier(criterion="entropy",n_estimators=10, random_state=0)
    rf.fit(x_train,y_train) y_rf_tr = rf.predict(x_train)
    print(accuracy_score(y_rf_tr,y_train))
    yPred_rf = rf.predict(x_test) = print(accuracy_score(yPred_rf,y_test))
    print("****Random Forest****")
    print("Confusion_Matrix")
    print(confusion_matrix(y_test,yPred_rf))
    print("Classification Report")
    print(classification_report(y_test,yPred_rf))

#saving the trained model
joblib.dump(rf, "RandomForestModel.pkl")

#Importing and building the KNN model
def KNN(x_train,x_test,y_train,y_test):
    knn = KNeighborsClassifier()
    knn.fit(x_train,y_train)
    y_knn_tr = knn.predict(x_train)
    print(accuracy_score(y_knn_tr,y_train))
    yPred_knn = knn.predict(x_test)
    print(accuracy_score(yPred_knn,y_test))
    print("****KNN****")
    print("Confusion_Matrix")
    print(confusion_matrix(y_test,yPred_knn))
    print("Classification Report")
    print(classification_report(y_test,yPred_knn))

#saving the trained model
joblib.dump(knn, "KNNModel.pkl")

#Importing and building the random forest model
def svm(x_tarin, x_test,y_train,y_test):
    svm = SVC(kernel = "linear")
    svm.fit(x_train,y_train)
    y_svm_tr = svm.predict(x_train)
    print(accuracy_score(y_svm_tr,y_train))
    yPred_svm = svm.predict(x_test)
    print(accuracy_score(yPred_svm,y_test))
    print("****Support Vector Machine****")
    print("Confusion_Matrix")
    print(confusion_matrix(y_test,yPred_svm))
    print("Classification Report")
    print(classification_report(y_test,yPred_svm))

#saving the trained model
joblib.dump(svm, "SVMModel.pkl")

# Importing the Keras libraries and packages
import keras
from keras.models import Sequential
from keras.layers import Dense

] # Initialising the ANN
classifier = Sequential()

# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 6, activation = 'relu', input_dim = 11))

# Adding the second hidden layer
classifier.add(Dense(units = 6, activation = 'relu'))

# Adding the output layer
classifier.add(Dense(units = 1, activation = 'sigmoid'))

# Compiling the ANN
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

# Training the ANN on the training set
classifier.fit(x_train, y_train, batch_size = 10, epochs = 100)

# Predicting the results
y_pred = classifier.predict(x_test)
y_pred = (y_pred > 0.5).astype('int')

# Making the Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)

# Saving the Model
model.save('ANN.h5')

```

```
[ ] # Adding the input layer and the first hidden layer
classifier.add(Dense (units=30, activation='relu', input_dim=40))

    func(x=x, v=v, **kwargs)

# Adding the second hidden layer
classifier.add(Dense (units=30, activation='relu'))

    func(x=x, v=v, **kwargs)

# Adding the output layer
classifier.add(Dense (units=1, activation='sigmoid'))

    func(x=x, v=v, **kwargs)

# Compiling the ANN
classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

    func(x=x, v=v, **kwargs)

#fitting the ANN to the training set
model_history=classifier.fit(x_train,y_train,batch_size=10,validation_split=0.33,epochs=200)

    func(x=x, v=v, **kwargs)

ann_pred = classifier.predict(x_test)
ann_pred = (ann_pred>0.5)
ann_pred

/usr/local/lib/python3.9/dist-packages/seaborn/axisgrid.py:1609: UserWarning: Ignoring `palette` because no `hue` variable has been assi
print(accuracy_score (ann_pred,y_test))
print("****ANN Model****")
print("Confusion_Matrix")
print(confusion_matrix(y_test, ann_pred))
print("Classification Report")

Saving... × nn_pred))

#testing on random input values
lr = LogisticRegression(random_state=0)
lr.fit(x_train,y_train)
print("Predicting on random input")
lr_pred_own = lr.predict(sc.transform([[0,0,1,1,0,0,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,1,0,0,456,1,0,3245,4567]]))
print("output is: ",lr_pred_own)

    func(x=x, v=v, **kwargs)

#testing on random input values
dtc = DecisionTreeClassifier (criterion="entropy", random_state=0)
dtc.fit(x_train,y_train)
print("Predicting on random input")
dtc_pred_own = dtc.predict(sc.transform([[0,0,1,1,0,0,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,1,0,0,456,1,0,3245,4567]]))
print("output is: ",dtc_pred_own)

    func(x=x, y=y, **kwargs)

#testing on random input values
rf = RandomForestClassifier (criterion="entropy",n_estimators=10, random_state=0)
rf.fit(x_train,y_train)
print("Predicting on random input")
rf_pred_own = rf.predict(sc.transform([[0,0,1,1,0,0,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,1,0,0,456,1,0,3245,4567]]))
print("output is: ",rf_pred_own)

/usr/local/lib/python3.9/dist-packages/seaborn/axisgrid.py:1609: UserWarning: Ignoring `palette` because no `hue` variable has been assi
#testing on random input values
Svc = SVC (kernel = "linear")
svc.fit(x_train,y_train)
print("Predicting on random input")
svm_pred_own = svc.predict(sc.transform([[0,0,1,1,0,0,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,1,0,0,456,1,0,3245,4567]]))
print("output is: ", svm_pred_own)

    func(x=x, y=y, **kwargs)

#testing on random input values
print("Predicting on random input")
ann_pred_own = classifier.predict(sc.transform([[0,0,1,1,0,0,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,0,0,1,1,0,0,456,1,0,3245,4567]]))
print (ann_pred_own)
ann_pred_own = (ann_pred_own>0.5)
print("output is: ", ann_pred_own)

    func(x=x, y=y, **kwargs)
```